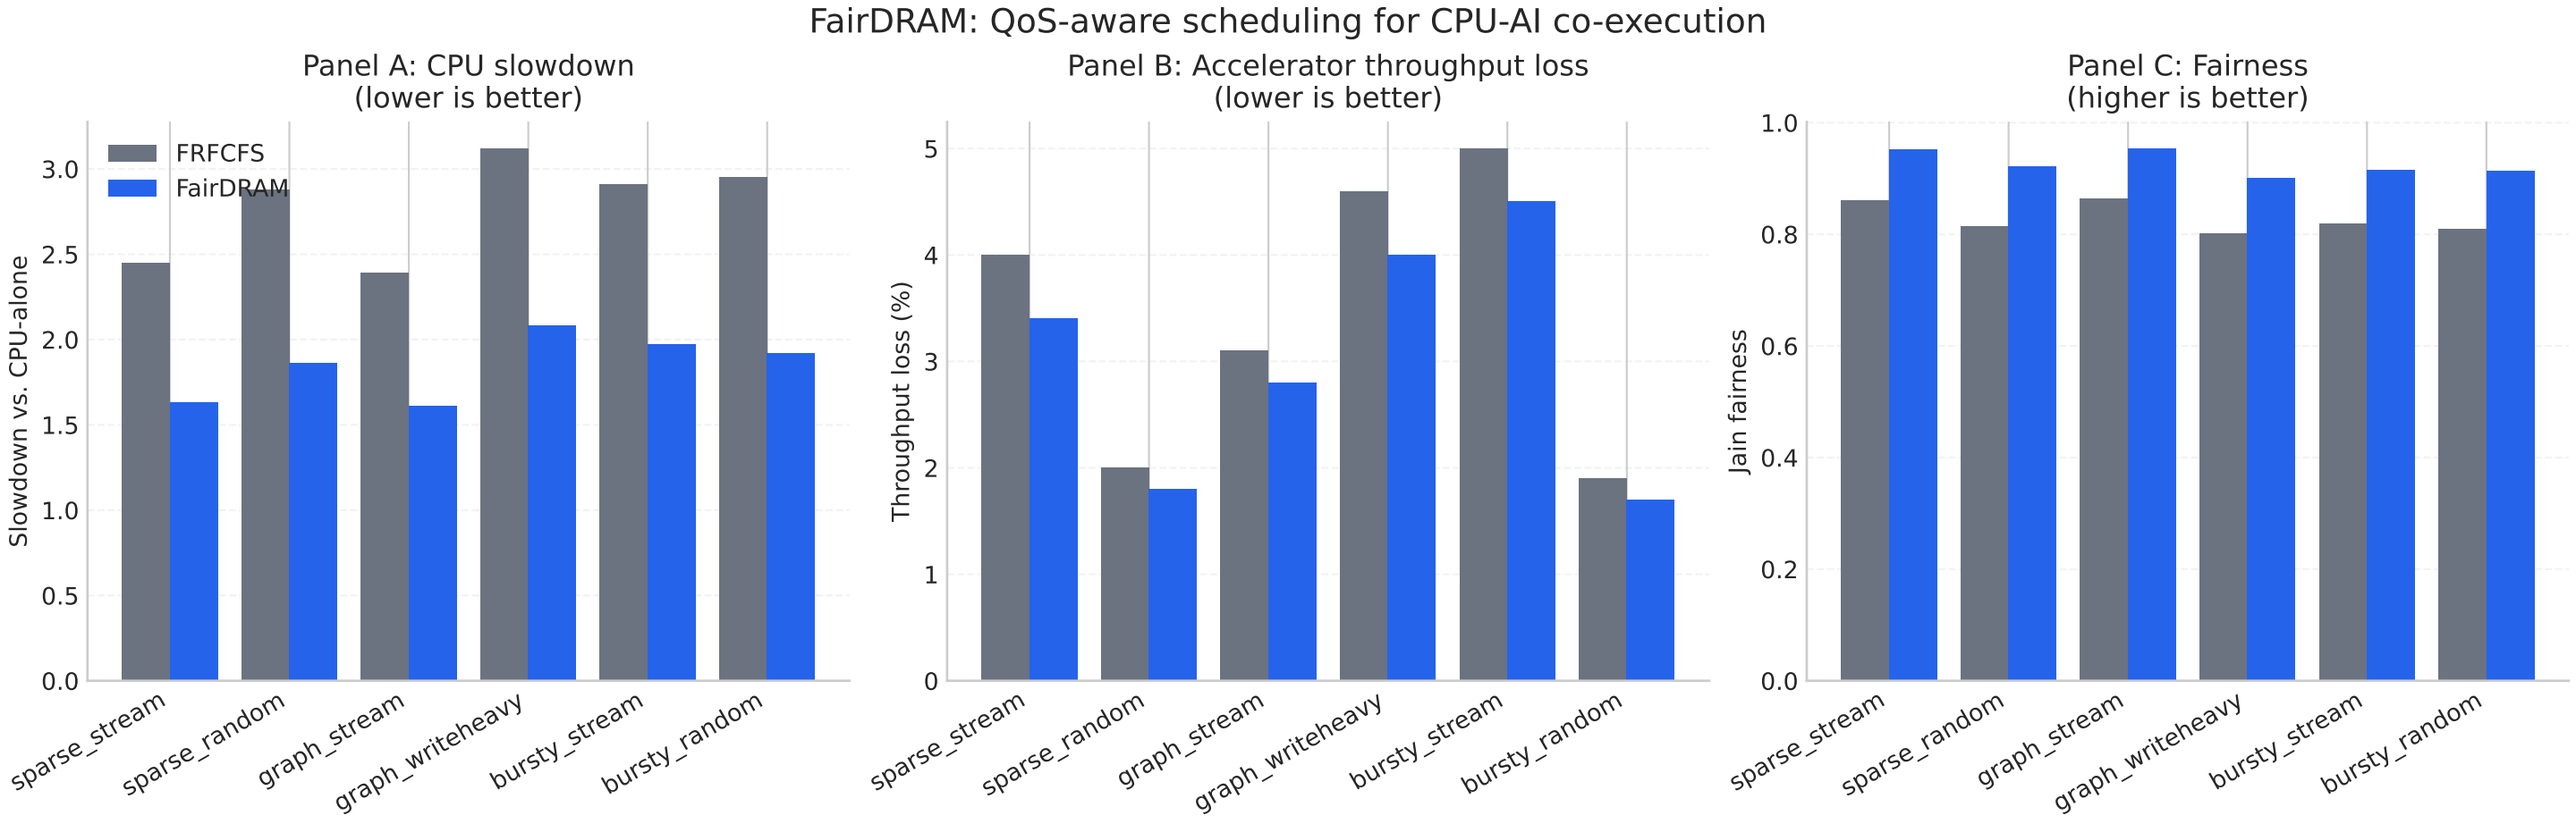
FairDRAM: QoS-aware scheduling for CPU-AI co-execution
Panel A: CPU slowdown
(lower is better)
Slowdown vs. CPU-alone
FRFCFS
FairDRAM
0.0
0.5
1.0
1.5
2.0
2.5
3.0
sparse_stream
sparse_random
graph_stream
graph_writeheavy
bursty_stream
bursty_random
Panel B: Accelerator throughput loss
(lower is better)
Throughput loss (%)
0
1
2
3
4
5
sparse_stream
sparse_random
graph_stream
graph_writeheavy
bursty_stream
bursty_random
Panel C: Fairness
(higher is better)
Jain fairness
0.0
0.2
0.4
0.6
0.8
1.0
sparse_stream
sparse_random
graph_stream
graph_writeheavy
bursty_stream
bursty_random